

Data exploration and analysis was completed utilizing big data tools such as HDFS, NiFi, Spark, and Solr.

Workout Activity

Week 11 - 12

Julie Campbell

The dataset that was chosen for the project consists of records of a person's physical activity recorded by a Redmi GPS Watch (Sharma, 2022). It includes data such as the date the activity was performed, workout type, duration, calories burned, and additional details about the result of the workout. Utilizing the tools discussed in the course, the data was explored and evaluated to determine how to maximize workout times to achieve fitness goals. For this project the workout duration, calorie burn, and heart rate will be the primary way to evaluate how to achieve the best workout pattern.

After downloading the data, I placed the file on the google virtual machine. NiFi was chosen as my first step to monitor and organize the flow of data between tools. Exploration was done on the 'InvokeHTTP' processor, but I found that it was more practical to replace the original file. I also did some exploration on 'GetHDFS' and 'FetchHDFS' based on the template that was provided through teams; however, I was unable to successfully connect to the HDFS instance even when inserting a temporary file. After all options were researched, I chose the 'GetFile' processor to modify and store the dataset. Once the file entered the flow, the 'QueryRecord' processor was used to remove the first column with the dummy ID. The final file replaces the original and the dataset was also transferred over to an existing Solr collection as shown in Figure 1.

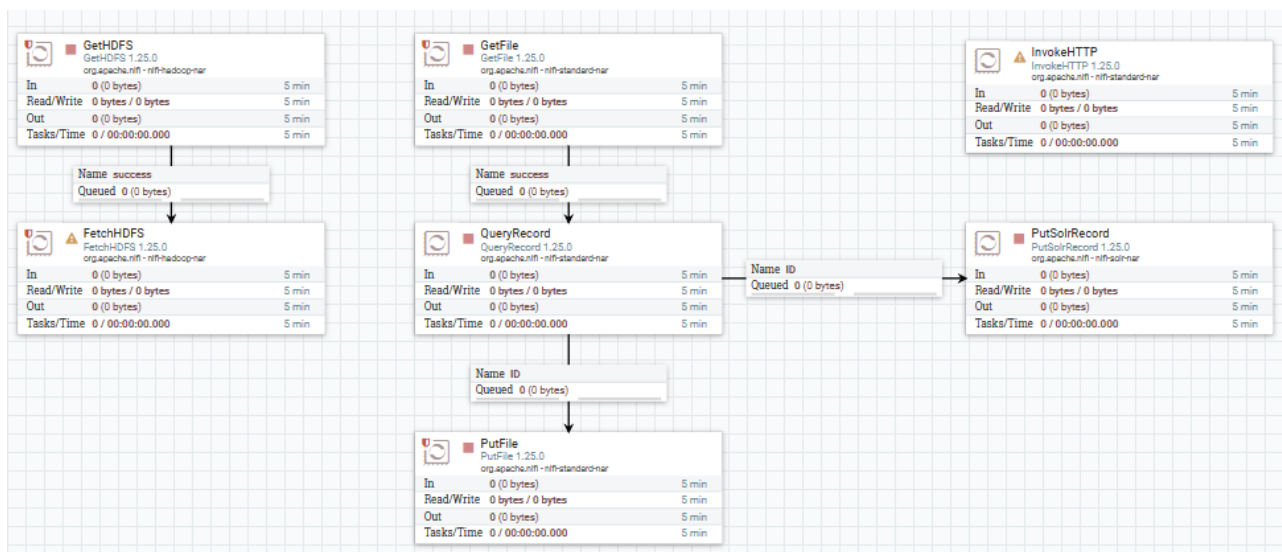


Figure 1: NiFi Flow

With the modified file in place, it was then loaded into HDFS to connect the dataset within the PySpark environment. During the initial load, I had noticed that the numeric data that I was utilizing was uploaded as strings and wouldn't allow me to perform the mathematical operations I needed. A schema was created utilizing the `pyspark.sql.types` library and struct types to define the format I wanted to load my data in as. The result of the schema change is shown in Figure 2.

```
>>> df.printSchema()
root
 |-- id: integer (nullable = true)
 |-- activity_day: date (nullable = true)
 |-- workout_type: string (nullable = true)
 |-- distance: double (nullable = true)
 |-- time: integer (nullable = true)
 |-- calories: integer (nullable = true)
 |-- total_steps: double (nullable = true)
 |-- avg_speed: double (nullable = true)
 |-- avg_cadence: double (nullable = true)
 |-- max_cadence: double (nullable = true)
 |-- avg_pace: string (nullable = true)
 |-- max_pace: string (nullable = true)
 |-- min_pace: string (nullable = true)
 |-- avg_heart_rate: double (nullable = true)
 |-- max_heart_rate: double (nullable = true)
 |-- min_heart_rate: integer (nullable = true)
 |-- vo2_max: integer (nullable = true)
 |-- aerobic: integer (nullable = true)
 |-- anaerobic: integer (nullable = true)
 |-- intensive: integer (nullable = true)
 |-- light: integer (nullable = true)
```

Figure 2 Pyspark Schema

I also imported the `pyspark.sql` functions library to perform operations such as mean, minimum, maximum, and count. Utilizing group by and aggregate, I was able to determine basic information about the workout, the heart rate range during workout, and the intensity category of the workout type. The results of the various queries performed are documented in Figure 3-5 with their respective lines of code. One of the biggest obstacles was that the results showed a mix of integers and floats. I had tried to lookup away to make the table reporting cleaner, but when I tried to change the data type in the schema it just removed the data since it didn't fit the pattern described (ex: float to int).

```
>>> df.groupby("workout_type").agg(F.count("workout_type"), F.mean("calories"),
,F.sum("distance"), F.sum("total_steps")).show()
```

workout_type	count(workout_type)	avg(calories)	sum(distance)	sum(total_steps)
Open Water	91	296.74725274725273	471.44000000000017	null
Walking	98	276.0408163265306	550.04000000000002	153958.0
Trail Run	90	267.96666666666664	463.32	251100.0
Freestyle	96	278.55208333333333	484.29999999999999	null
Trekking	94	283.1276595744681	496.19	472632.0
Indoor Cycling	80	280.45	420.72000000000014	null
Outdoor Running	81	301.4691358024691	408.88999999999993	617058.0
Cricket	93	307.5483870967742	491.39000000000016	320292.0
Treadmill	98	278.14285714285717	489.98999999999995	445900.0
Pool Swimming	94	283.4148936170213	504.6	null
Outdoor Cycling	85	299.1294117647059	462.73999999999999	null

Figure 3 Workout Summary

The results of the query in Figure 3 show that walking and treadmill were the most utilized workout types, but the top calorie burning workouts were cricket and outdoor running.

```
>>> df.groupby("workout_type").agg(F.min("min_heart_rate"), F.mean("avg_heart_rate"),
,F.max("max_heart_rate")).show()
```

workout_type	min(min_heart_rate)	avg(avg_heart_rate)	max(max_heart_rate)
Open Water	80	130.7032967032967	166.0
Walking	80	90.31122448979592	85.0
Trail Run	81	141.98333333333332	188.0
Freestyle	80	109.44791666666667	122.0
Trekking	80	110.35106382978724	124.0
Indoor Cycling	80	92.4	90.0
Outdoor Running	80	143.57407407407408	190.0
Cricket	80	93.51075268817205	92.0
Treadmill	80	100.29591836734694	105.0
Pool Swimming	80	135.41489361702128	175.0
Outdoor Cycling	80	136.50588235294117	177.0

Figure 4 Workout Heart Rate

In Figure 4, I compared the heart rate range between the various workout types. The results showed anomalies with walking, indoor cycling, and cricket. After further analysis, I discovered that the data captured for max heart rate was entered as one number across all records in the anomalies. For the purposes of analysis, this will carry less weight than other features.

```
>>> df.groupBy("workout_type").agg(F.mean("aerobic"),F.mean("anaerobic"),F.mean("intensive"),F.mean("light")).show()
```

workout_type	avg(aerobic)	avg(anaerobic)	avg(intensive)	avg(light)
Open Water	0.0	5.0	38.15384615384615	50.252747252747255
Walking	22.0	0.0	34.785714285714285	45.11224489795919
Trail Run	0.0	80.0	38.08888888888889	50.98888888888889
Freestyle	28.0	2.0	37.635416666666664	49.5625
Trekking	50.0	0.0	36.02127659574468	53.329787234042556
Indoor Cycling	32.0	0.0	39.6	51.1
Outdoor Running	0.0	71.0	32.888888888888886	55.160493827160494
Cricket	40.0	10.0	36.924731182795696	58.8494623655914
Treadmill	26.0	7.0	37.826530612244895	45.87755102040816
Pool Swimming	20.0	15.0	33.08510638297872	51.93617021276596
Outdoor Cycling	45.0	15.0	33.470588235294116	49.247058823529414

Figure 5 Workout Intensity

The results from Figure 5 show an average of the category percentage that each workout type had assigned per session. Instead of utilizing average, I wanted to apply a ranking system to see which category best fits the workout type to get a baseline for the intensity. I tried out a method using window and rank as suggested by a stack overflow post, but I was unsuccessful (Group By, Rank and aggregate spark data frame using pyspark, n.d.) as shown in Figure 6.

```
df.select("workout_type","light").withColumn("dense_rank",dense_rank().over(windowSpec)).show()
```

Figure 6 Workout Intensity Rank

The final tool that was used during analysis was Solr. According to heart.org, a healthy 35-year-old should target a heart rate of 93-157 bpm with a maximum of 185bpm (Target Heart Rates Chart, 2021). This reference will be used as a baseline to evaluate the person's workout routine. The average and max heart rates were filtered to meet the criteria. The top workout descriptions that met the criteria were cycling, freestyle, pool, swimming, and treadmill as shown in Figure 7.

```
{
  "facet_counts":{
    "facet_queries":{ },
    "facet_fields":{
      "workout_type":["cycling",122,"freestyle",96,"pool",94,"swimming",94,"treadmill",94,"trekking",94,"open",91,"water",91,"outdoor",85]
    },
    "facet_ranges":{ },
    "facet_intervals":{ },
    "facet_heatmaps":{ }
  }
}
```

Figure 7 Workout Facet Counts

The top 10 calorie burn entries were also recorded with the filters for target and max heart rate range. The top workout types were open water, trekking, indoor cycling, freestyle, outdoor cycling, and treadmill. This is shown in Figure 8. Compared to the frequent workout types, there appears to be diversity in the types that will achieve the results wanted.

The screenshot displays the Solr Request-Handler (qt) interface on the left and the resulting JSON response on the right.

Request-Handler (qt) Interface:

- URL:** /select
- q:** avg_heart_rate:[93 TO 157]
max_heart_rate:[80 TO 185]
- q.op:** AND
- fq:** (empty)
- sort:** calories desc, activity_day desc
- start, rows:** 0, 10
- fl:** workout_type, calories
- df:** (empty)
- paramset(s):** Select paramset(s)...
- wt:** json
- indent on:** ☒
- debugQuery:** ☐
- defType:** -----
- hl:** ☐
- facet:** ☒
- facet.query:** (empty)
- facet.field:** workout_type
- facet.prefix:** (empty)
- facet.contains:** (empty)
- facet.contains.ignoreCase:** ☐
- facet.limit:** (empty)
- facet.matches:** (empty)
- facet.sort:** count

JSON Response:

```
{
  "responseHeader": {
    "status": 0,
    "QTime": 3,
    "params": {
      "facet.field": "workout_type",
      "indent": "true",
      "fl": "workout_type, calories",
      "q.op": "AND",
      "sort": "calories desc, activity_day desc",
      "rows": "10",
      "q": "avg_heart_rate:[93 TO 157]\\nmax_heart_rate:[80 TO 185]",
      "facet.mincount": "50",
      "facet.contains.ignoreCase": "false",
      "wt": "json",
      "facet.sort": "count",
      "useParams": "",
      "_: "1709160814848"
    }
  },
  "response": {
    "numFound": 673,
    "start": 0,
    "numFoundExact": true,
    "docs": [
      {
        "workout_type": ["Open Water"],
        "calories": [550]
      },
      {
        "workout_type": ["Trekking"],
        "calories": [550]
      },
      {
        "workout_type": ["Trekking"],
        "calories": [550]
      },
      {
        "workout_type": ["Treadmill"],
        "calories": [549]
      },
      {
        "workout_type": ["Indoor Cycling"],
        "calories": [548]
      },
      {
        "workout_type": ["Freestyle"],
        "calories": [547]
      },
      {
        "workout_type": ["Outdoor Cycling"],
        "calories": [545]
      },
      {
        "workout_type": ["Open Water"],
        "calories": [545]
      },
      {
        "workout_type": ["Freestyle"],
        "calories": [545]
      },
      {
        "workout_type": ["Treadmill"],
        "calories": [544]
      }
    ]
  }
}
```

Figure 8 Workout Solr Results

In conclusion, it appears that the person's workout routines should primarily focus on cycling, freestyle, pool swimming, treadmill, and trekking. The big data tools (HDFS, NiFi, Spark, and Solr) provided the framework to explore the data and organize the flow. Next steps would be to build a machine learning model to provide a personalized workout recommendation for multiple users.

References

- Group By, Rank and aggregate spark data frame using pyspark*. (n.d.). Retrieved from Stack Overflow: <https://stackoverflow.com/questions/41661068/group-by-rank-and-aggregate-spark-data-frame-using-pyspark>
- Sharma, T. (2022). *Redmi Fuel Band Record Tracker (Fitbit Dataset)*. Retrieved from Kaggle: <https://www.kaggle.com/datasets/tanisha1416/my-redmi-fuel-band-record-tracker-fitbit-dataset?rvi=1>
- Target Heart Rates Chart*. (2021). Retrieved from Heart: <https://www.heart.org/en/healthy-living/fitness/fitness-basics/target-heart-rates>